Venkata Lakshmi Parimala Pasupuleti

venkatalakshmiparimala@gmail.com | +1-703-261-5410 | Virginia, USA

September 21, 2026

Hiring Manager, AI Engineering
FourKites
Chicago, Illinois

I'm applying for the Senior AI Engineer role on FourKites' AI agent team. Your job description mentions a support automation agent handling 60-70% of customer tickets and an end-to-end carrier onboarding agent. I've built exactly that - not as a concept, but in production code I can walk through line by line.

Leasel is an 8-agent LangGraph support investigation platform: Triage → Retrieval → Configuration → Historical Incident → Structured Data → Reasoning → Validation → Human Escalation, with conditional routing and HITL escalation when confidence drops below threshold. It classifies support tickets, retrieves relevant business rules and historical incidents, and produces ranked, evidence-backed root-cause hypotheses - the same architecture your team is deploying for carrier support automation. Evaluated across 1,240 test cases: 93.6% precision, 87.9% F1. **AgentProbe** is the evaluation toolkit I built to benchmark it: 5-metric eval engine (correctness, groundedness, hallucination rate, latency, token efficiency), LLM-as-judge scoring, CI-integrated regression gate that blocks quality degradation on every commit, and multi-provider LLM routing across Anthropic, OpenAI, Groq, and Ollama. When FourKites needs to measure whether your support agent is getting better or worse, that's the kind of infrastructure I build to answer that question.

Before those projects, I spent 7 years at Entrata shipping backend systems in Python, Java, and PostgreSQL across 120+ Agile sprints - owning production modules end-to-end through development, QA, and release. RabbitMQ async pipelines, microservices redesign in Docker (50% scalability improvement), 200+ SQL queries optimized, 300+ code reviews conducted. The engineering rigor that makes AI agents actually reliable in production - that's where that foundation comes from.

FourKites processes 3.2 million supply chain events daily. Building agents that operate reliably at that volume, with the evaluation infrastructure to prove they're working - that's exactly the problem I want to be solving. I'd welcome the chance to discuss how my background maps to what your team is building.

Sincerely,

Venkata Lakshmi Parimala Pasupuleti