Venkata Lakshmi Parimala Pasupuleti

venkatalakshmiparimala@gmail.com | +1-703-261-5410 | Virginia, USA

September 21, 2026

Hiring Manager, Applied AI Engineering
LVT
American Fork, Utah

I'm applying for the Senior Software Engineer, Applied AI role at LVT. Your job description asks for "a product-level evaluation and regression suite that catches false or mistimed talkdowns before they reach the field." I built that - not as a concept, but as a production system I can walk through commit by commit.

AgentProbe is a CI-integrated AI evaluation and regression framework: 5-metric eval engine (correctness, groundedness, hallucination rate, latency, token efficiency), LLM-as-judge scoring with multi-trial pass@k harness, and a regression gate that blocks quality degradation on every commit before anything ships. It runs across 1,240+ test cases and gives a traceable rationale for every score. That is the tooling a real-world-actuating system needs - the thing that answers "did it act correctly this time, and will it act correctly next time" before the answer matters in the physical world. **Leasel** is the agentic platform it was built to evaluate: an 8-agent LangGraph orchestration system with prompting, RAG retrieval, tool calling, conditional routing, and HITL escalation - where the core design rule enforced in code is: LLM reasons, deterministic logic makes every final call. 93.6% precision across 1,240 compliance test cases.

The precision discipline your role requires is one I know from a different domain: 7 years at Entrata owning systems where a wrong data value determines housing eligibility for thousands of residents and triggers federal regulatory corrections. Zero compliance violations across the full ownership period - not because the stakes were low, but because the engineering practice made wrong outputs structurally unlikely and detectable before they propagated. That same discipline - instrument everything, gate on quality, own the bar - is what a detection-to-response system acting in the physical world demands.

I'm genuinely excited about what LVT is building. The intersection of physical-world detection and LLM orchestration - owning both the GenAI harness and the evaluation infrastructure that makes it trustworthy - is exactly the problem space I want to work in. I'd welcome the opportunity to discuss how my background maps to this role.

Sincerely,

Venkata Lakshmi Parimala Pasupuleti