Venkata Lakshmi Parimala Pasupuleti

Software Engineer · Production Systems · Performance Engineering · Security · ML Integration · CI/CD

venkatalakshmiparimala@gmail.com | +1-703-261-5410 | linkedin.com/in/venkata-lakshmi-parimala-pasupuleti | github.com/venkatalakshmiparimala

Software Engineer — Face ID · Apple

September 24, 2026

Dear Hiring Team,

Face ID runs billions of times a day on hardware that can't afford a wrong answer — wrong in latency, wrong in power, wrong in security, wrong in the authentication decision itself. The engineering discipline that requires is what drew me to this role. I want to build software where correctness is not optional and shipping is the proof.

My production engineering background is nine years deep. At Entrata I owned backend systems serving 100K+ monthly active users at 99.5% uptime — profiled and optimized 200+ queries for a 30% performance gain, rebuilt microservices in Docker for a 50% scalability improvement, and instrumented the stack end-to-end with OpenTelemetry, Prometheus, and Grafana (40% reduction in mean time-to-diagnose). Every production incident I owned ended with a permanent root-cause fix, not a workaround. On security: I led the AWS S3 migration for Affordable Housing document workflows — Libsodium-encrypted pipelines across hundreds of thousands of PHI records, HIPAA-compliant RBAC, zero security incidents across the entire ownership period. That combination of reliability, performance rigor, and security discipline maps directly to what the Face ID stack demands.

On the ML integration side: I've shipped AI components into production APIs — BERT embeddings and LLMs into a HIPAA-compliant FHIR medical records system, 8-agent LangGraph workflows into a property operations platform with deterministic safety boundaries enforced in code, and a full CI-integrated benchmarking harness (AgentProbe) that gates every model change on quantitative regression before it ships. I work alongside the models; I don't need to invent them. I understand the integration contract, the failure modes, and what it takes to make an ML component reliable enough to run on a device that can't phone home for a second opinion.

C/C++ and the Apple platform are the steepest part of the learning curve here and I won't pretend otherwise. What I bring is nine years of production engineering discipline, a track record of owning systems under real constraints, and the demonstrated ability to ramp fast. I'd welcome the conversation.

Sincerely,

Venkata Lakshmi Parimala Pasupuleti

venkatalakshmiparimala@gmail.com · +1-703-261-5410 · github.com/venkatalakshmiparimala