Venkata Lakshmi Parimala Pasupuleti

venkatalakshmiparimala@gmail.com | +1-703-261-5410 | github.com/venkatalakshmiparimala

September 24, 2026

Dear LlamaIndex Engineering Team,

I've been building the things LlamaIndex makes possible. LeaseI, my document understanding platform, does OCR-fallback ingest, section-aware chunking, hybrid search (TF-IDF + semantic reranking), and citation grounding to exact source clauses - the same building blocks LlamaIndex ships as infrastructure. AgentProbe, my open-source AI eval toolkit, went from prototype to production because the tooling didn't exist and the problem needed solving. Both of those decisions - identifying what the ecosystem needs, building it anyway - are exactly what I see in the engineering philosophy behind LlamaIndex.

My backend foundations match your stack directly: Python (8+ years), PostgreSQL, Kubernetes, Terraform, Temporal, RabbitMQ, and AWS. At Entrata, I owned 20+ production modules serving 100K+ monthly active users at 99.5% uptime across 120+ Agile sprints - including a mission-critical compliance pipeline I operated solo for 7 years without a project manager or a playbook. I know what it means to own something end-to-end: design, build, ship, operate, and fix when it breaks at 2am.

What draws me specifically to LlamaIndex is the product surface. Document understanding and knowledge retrieval aren't adjacent to the work I've been doing - they are the work. LeaseI processes residential lease PDFs across 5 US states and federal regulation, achieves 93.6% precision across 1,240 test cases, and enforces a hard rule: LLM reasons, deterministic logic decides every regulated outcome. That's the kind of production thinking I'd bring to scaling out LlamaIndex's ingestion and AI systems.

I'd welcome the chance to talk about how my experience building on top of - and building the infrastructure behind - RAG and document understanding systems maps to what your team is scaling next.

Sincerely,

Venkata Lakshmi Parimala Pasupuleti