Venkata Lakshmi Parimala Pasupuleti

+1-703-261-5410 | venkatalakshmiparimala@gmail.com | github.com/venkatalakshmiparimala | Virginia, USA

October 2, 2026

Hiring Team
Cardiff
Remote

I'm applying for the Sr. AI Product Engineer role. Cardiff's goal - reducing application-to-funding from 24 hours to minutes - is exactly the kind of problem I want to work on, and the way you've described this role maps precisely to how I work.

The job is assess, harden, ship, own. I've done this. **Leasel** (github.com/venkatalakshmiparimala/Lease_complaince_rag) started as a prototype idea. I reviewed the architecture options, documented what should move forward and what shouldn't, built the data models, APIs, error handling, tests, and deployment process - then shipped three production surfaces that operate today with a CI regression gate and monitoring. That's the 30/90/180-day arc you've laid out, end to end.

On the "carefully review AI-generated code" requirement: I built **AgentProbe** (github.com/venkatalakshmiparimala/AgentProbe) specifically because I didn't trust AI-generated outputs without a harness. It's an open-source LLM evaluation toolkit with CI-integrated regression testing, automated golden-dataset validation, and LLM-as-judge scoring. I use it as a release control layer before anything ships. That's the mindset Cardiff is asking for.

Before these projects, I spent nine years at Entrata owning enterprise system integrations, data pipelines, and production incident response on a regulated platform serving 100K+ users. I know what it takes to keep a system running after it ships - not just get it out the door. Your stack (AWS, Salesforce, Snowflake) is familiar ground.

I'd like to talk. Happy to walk through either project or discuss how I'd approach the prototype portfolio assessment in the first 30 days.

Sincerely,

Venkata Lakshmi Parimala Pasupuleti