Venkata Lakshmi Parimala Pasupuleti

Builder · AI-First · Zero-to-One · Shipped Projects · LangGraph · RAG · LLM Evaluation · Python · AWS

+1-703-261-5410 | venkatalakshmiparimala@gmail.com | linkedin.com/in/venkata-lakshmi-parimala-pasupuleti | github.com/venkatalakshmiparimala | Virginia, USA

SUMMARY

Engineer who builds things outside of work and finishes them. Built **Leasel** independently - from blank page to production: 8-agent LangGraph system with RAG pipelines, hybrid search, HITL escalation, and a CI evaluation harness. 93.6% precision across 1,240 test cases. Open source, in production, real users. Built **AgentProbe** - open-source LLM evaluation toolkit, multi-provider routing, CI regression harness - because I needed it and it didn't exist the way I wanted it. Currently at Intuit integrating Anthropic Claude into production financial services. Before that, 7+ years at Entrata owning production systems at scale - the kind of work that builds the instinct for when to go fast, when to be careful, and how to keep a system alive after it ships. M.S. Data Analytics Engineering, George Mason University (GPA 3.87). Comfortable picking up new tools fast. Have strong opinions about what should be built and why, and I'll argue for them.

CORE COMPETENCIES

Zero-to-One Builder | AI-First Product Development | Prototyping · Shipping · Iteration | Agentic AI · LangGraph · LangChain | RAG Pipelines · Hybrid Search · Embeddings | LLM Evaluation · Evals · CI Regression | Anthropic Claude · OpenAI · LLM APIs | Python · FastAPI · TypeScript · React | AWS · Docker · GitHub Actions | Self-Driven Learning · Frontier AI Tools | Production Engineering · System Design | Open Source · Side Projects That Shipped

EXPERIENCE

Software Engineer
Jun 2026 - Present

Intuit — Virginia, USA

- Integrating Anthropic Claude into production Python/FastAPI financial services - structured outputs, retrieval grounding, prompt design, LLM-as-judge eval regression tests, and guardrails. Shipping and iterating in production with Docker/GitHub Actions CI/CD to AWS.
- Building ML pipelines (scikit-learn, XGBoost) and React/TypeScript front ends; making architectural decisions about what to build and how to measure it.

Senior Software Engineer
Oct 2020 - Jul 2024

Entrata — Apartment & Affordable Housing SaaS (100K+ monthly active users)

- Owned production systems end-to-end for 7+ years - compliance data pipelines, async processing, AWS infrastructure, microservices - in a regulated multi-tenant platform where wrong outputs have real consequences for real people.
- Delivered 20+ features across 120+ sprints; Python/Bash automation cutting 80% of manual work; mentored junior engineers; 300+ code reviews; built internal debugging tools that cut triage time by 40%.

Junior Software Engineer
Sep 2018 - Oct 2020

Entrata India — Pune

- Built compliance and property management modules (Python, PHP, Java, PostgreSQL); led AWS S3 migration; delivered large-scale data import workflows (1.6M records).

SHIPPED PROJECTS

Leasel - Built from Scratch, In Production, Open Source
2026 | Python, LangGraph, FastAPI, ChromaDB, CI/CD | github.com/venkatalakshmiparimala/Lease_complaince_rag

- Saw a real problem, framed the solution, built the whole thing - 8-agent LangGraph system with tool calling, durable state, RAG pipeline (OCR ingestion, section-aware chunking, hybrid search, semantic reranking), and a CI evaluation harness with golden datasets and LLM-as-judge scoring. Hard rule enforced in code: AI reasons, human or deterministic logic makes the compliance call. 93.6% precision across 1,240 test cases. Didn't stop at prototype - shipped it, instrumented it, and operate it.

AgentProbe - Built Because It Needed to Exist
2025 | Python, FastAPI, LangGraph, CI/CD | github.com/venkatalakshmiparimala/AgentProbe

- Open-source LLM evaluation toolkit: offline evals, golden datasets, LLM-as-judge scoring, CI regression harness, multi-provider routing across Anthropic, OpenAI, Azure OpenAI, Groq, and Ollama. Built it because I needed a proper eval harness for Leasel and the existing tools didn't do what I wanted. Used daily as a release gate before AI systems ship to production.

AI on FHIR - Clinical NLP App
2024 | Python, Flask, React, BERT, LLMs, FHIR | github.com/venkatalakshmiparimala/AI_on_FHIR

- Translates English/Spanish natural-language questions into FHIR API calls using LLMs, BERT, and embeddings - cut manual clinical lookup time ~60%. HIPAA-style RBAC, containerized with Docker.

EDUCATION

M.S. Data Analytics Engineering — George Mason University, Virginia | GPA 3.87 — May 2026

B.Tech Computer Science & Engineering — JNTU Anantapur — 2017