Venkata Lakshmi Parimala Pasupuleti

lakshmiparimalap@gmail.com | +1-703-261-5410 | linkedin.com/in/venkata-lakshmi-parimala-pasupuleti | Virginia, USA

Dear Lazarus AI Hiring Team,

Your description of the problem - AI initiatives stalling because teams can't get from experimentation to production safely - is exactly what I've spent the last two years solving in code.

AgentProbe, my open-source LLM evaluation platform, is a CI-integrated quality gate: 5-metric eval engine (correctness, groundedness, hallucination rate, latency, token efficiency), LLM-as-judge scoring, and a regression suite that blocks any prompt change that degrades performance before it reaches production. Zero regressions shipped across 1,240+ test cases. The infrastructure is built before any prompt touches users - evaluation is the on-ramp to production, not an afterthought.

LeaseI applies the same philosophy to a regulated domain - affordable housing compliance. Every output is either cited to an exact source clause or flagged as unsupported. Every eligibility and financial decision routes to human review. The design rule is enforced in code, not policy: the LLM extracts and reasons; deterministic or human logic makes the final regulated call. Precision: 93.6% across 1,240 compliance test cases. This is what "auditable and explainable" looks like in practice. The regulated-domain grounding isn't borrowed - I spent 7 years at Entrata owning the HUD income-limit pipeline that powered eligibility decisions for 100K+ housing units, with 100% PHI compliance and zero security incidents across the entire period.

Lazarus AI is building infrastructure to make production AI reliable and trustworthy in regulated environments. That is the exact problem I've been building toward. I'd welcome the chance to talk through how my work fits what you're building.

Sincerely,

Venkata Lakshmi Parimala Pasupuleti

github.com/venkatalakshmiparimala/AgentProbe