Venkata Lakshmi Parimala Pasupuleti

venkatalakshmiparimalap@gmail.com | +1-703-261-5410 | Fairfax, VA | LinkedIn | GitHub

SUMMARY

AI practitioner and software engineer with 7+ years in **multifamily real estate technology** (Entrata -- same industry as TheGuarantors: owners, operators, affordable housing) and deep hands-on experience evaluating, testing, and deploying AI solutions. Built **AgentProbe**, a multi-provider LLM evaluation platform that benchmarks AI outputs across five vendors on accuracy, hallucination rate, and cost -- directly mirroring the vendor research and pilot testing this role requires. Strong **analytical skills** -- Python and SQL for operational data analysis, clear communication of technical findings in writing and presentations, and translating AI capabilities into measurable business outcomes. M.S. Data Analytics Engineering, George Mason University (GPA 3.87, May 2026).

CORE COMPETENCIES

AI Vendor Evaluation & Pilot Testing | Third-party AI Platform Research | Python / SQL / Operational Data Analysis | LLM Tools (Claude, GPT, Llama) | RAG Pipelines / AI Orchestration | AI Performance Metrics & Benchmarking | Workflow Analysis & Documentation | Workflow Automation | Software Testing & QA | Technology Implementation | Stakeholder Communication | Business Operations & Operational Improvement | Cost-Benefit Analysis for AI Initiatives | Multifamily Real Estate & Leasing Risk Domain | REST APIs / Backend Integration

EXPERIENCE

Intuit

Jun 2026 -- Present

Software Engineer -- AI/LLM Tools & Production Systems -- Virginia, USA

- **AI initiatives and third-party platform evaluation:** research AI tools and vendors, compare capabilities and pricing, summarize findings and open questions, and recommend guardrails and governance standards for production adoption.
- Support AI pilot testing: define test scenarios, check outputs against agreed criteria, document issues and user feedback, and prepare draft updates with suggested next steps for stakeholder review.
- Build and run **evaluation harnesses** (LLM-as-judge, golden datasets) to track agreed performance metrics -- model accuracy, hallucination rates, completeness -- before enterprise deployment.
- Analyze operational bottlenecks to size opportunities, estimate potential benefits and costs, and determine whether workflow automation, rule-based logic, or generative AI is the right approach.
- Build ML pipelines (scikit-learn, XGBoost) and deploy AI tools via Docker and GitHub Actions CI/CD on AWS.

Entrata -- Multifamily & Affordable Housing SaaS

Oct 2020 -- Jul 2024

Senior Software Engineer -- Multi-tenant platform, 100K+ monthly active users

- Operated in the **same multifamily ecosystem** as TheGuarantors -- built software used daily by property owners, operators, and leasing teams to manage resident eligibility, affordable housing compliance, and rent processing.
- Analyzed operational data and workflows across multiple business teams to scope features, estimate impact, and document requirements -- working with stakeholders from operations, compliance, and leasing.
- Optimized 200+ SQL queries via execution plan analysis (30% performance gain); built Python/Bash automation eliminating 80% of manual operational workload.
- Designed Docker microservices on AWS with Terraform; instrumented for production observability; reviewed 300+ PRs and mentored junior engineers.

Entrata India

Sep 2018 -- Oct 2020

Junior Software Engineer -- Pune, India

- Built property management modules (Python, PostgreSQL) for affordable housing and rental management; delivered large-scale data import workflows (1.6M HUD/Tax Credit records).

PROJECTS

AgentProbe -- AI Platform Evaluation & Vendor Benchmarking Tool

Python · FastAPI · 5 LLM Providers · Cost Analysis · CI/CD

Built a 15-endpoint evaluation platform that routes prompts across five AI providers (Anthropic, OpenAI, Azure OpenAI, Groq, Ollama) and scores outputs on factual accuracy, hallucination rate, retrieval quality, completeness, and compliance -- with cost-aware routing. Directly mirrors the vendor research, capability comparison, and pilot testing this role performs. CI regression against golden datasets tracks performance over time.

github.com/venkatalakshmiparimala/AgentProbe

Leasel -- AI Pilot: Lease Compliance with Measured Outcomes

Python · LangGraph · FastAPI · RAG · Eval Harness · 1,240 Test Cases

End-to-end AI pilot for apartment lease compliance: defined test scenarios, built the RAG pipeline, measured output quality with a structured evaluation harness, and documented precision/F1 results across 1,240 cases (93.6% precision, 87.9% F1). Human escalation at defined confidence thresholds. Demonstrates the full pilot lifecycle -- design, test, measure, document -- that this role supports.

github.com/venkatalakshmiparimala/Lease_complaince_rag

AI on FHIR -- Clinical Workflow Automation

Python · Flask · React · FHIR · RBAC · Docker

Evaluated AI feasibility for clinical query workflows, built a prototype translating English/Spanish queries into FHIR API calls using LLMs and BERT, and demonstrated ~60% reduction in manual lookup time. HIPAA-style RBAC and containerized deployment.

github.com/venkatalakshmiparimala/AI_on_FHIR

EDUCATION

M.S. Data Analytics Engineering — George Mason University, Virginia — May 2026 | GPA 3.87

B.Tech Computer Science & Engineering — JNTU Anantapur — 2017